

Guardrail Module: Prompt Injection, Jailbreak, Exfiltration, Toxicity, and PII Detection

By Insharah Nazir

1. Executive Summary.....	1
2. Architecture.....	2
2.1 Interception Points.....	2
2.2 Detection Modules.....	2
3. Tool Selection: What Was Tested and Rejected.....	3
4. Attack Prompt Test Set.....	4
5. Evaluation Results.....	5
5.1 Fast Layer vs. Fast+LLM: Overall Comparison.....	6
5.2 By Category.....	7
5.3 By Language.....	8
6. Design Decisions Made From Evidence.....	9
7. Known Limitations.....	9
8. Conclusion.....	10

1. Executive Summary

This report documents the design, implementation, and evaluation of a multilingual (English, Korean, Urdu) guardrail module for the RAG system built in Phase 3. The module detects five categories of adversarial or unsafe input: prompt injection, jailbreak attempts, data exfiltration, toxic content, and PII leakage.

The guardrail was built and evaluated in two stages. A fast layer — a small pretrained classifier plus hand-built regex, no LLM calls — was built first and evaluated against a 125-prompt labeled test set (94 attacks, 31 benign controls, spanning all three languages), reaching 84.0% accuracy and 98.7% precision. Investigation of the fast layer's remaining misses showed they concentrated almost entirely in narrative/roleplay-framed jailbreak attempts — a category that pattern-matching tools cannot reliably resolve regardless of how much they are tuned, since the attack is expressed through structure and implication rather than a surface-level trigger phrase.

An LLM-based semantic judge was added as a second-pass check, reached only when the fast layer finds nothing (so most traffic never incurs the added cost). This is the configuration adopted for the production system. Evaluated against the same 125-prompt set, the fast+LLM guardrail reaches **92.8% accuracy, 91.5% recall, and 98.9% precision**, with an attack success rate of **8.5%**. This is the headline result of this report; the fast-layer-only numbers are retained throughout as a comparison baseline to show what the added LLM stage specifically bought.

2. Architecture

2.1 Interception Points

The guardrail runs as two interception points around the existing RAG pipeline:

- **check_input_deep(text, language)** — the production entry point. Runs the fast layer (check_input) first; if it finds nothing, a single LLM call (judge_query) checks the query for narrative/roleplay jailbreak intent before the query is allowed to proceed to retrieval. If the fast layer already blocks, the LLM call is skipped entirely — most traffic (roughly 91% in evaluation) resolves at the fast layer alone, keeping the average added cost low.
- **check_output(text, language)** — runs on the generated answer before it is returned to the user. Screens for exfiltration (in case retrieved corpus content successfully injected something into the answer) and PII leakage. This stage does not include the LLM judge, since its purpose (catching narrative jailbreak intent in a *query*) does not apply to a generated *answer*.

Both the fast layer and check_output run once per request, not once per query variant or retrieved chunk. The LLM judge, when it fires, adds one additional model call on top of the existing rewrite/multi-query/citation calls already in the pipeline.

2.2 Detection Modules

Injection / direct jailbreak: Meta Llama Prompt Guard 2 (86M parameters). A binary classifier fine-tuned to detect prompt injection and direct jailbreak framing. Selected after head-to-head testing against ProtectAI's DeBERTa-v3 model (see Section 3). Runs as the first check in the fast layer; if it flags a query above threshold, the request is blocked immediately and no further checks or the LLM judge are reached.

Jailbreak roleplay/persona-override: custom regex (jailbreak_detector.py), EN/KO/UR. Targets a specific, narrower pattern than the LLM judge below: explicit persona-override phrasing ("pretend you have no restrictions," "you are now DAN," "act as an AI with no rules"). This was written to close a gap found during smoke testing of Prompt Guard 2, which caught this pattern reliably in English but missed it in Korean and Urdu. It is intentionally narrow — a fast, deterministic first pass for the most common, templated form of this attack, distinct from the LLM judge's broader semantic coverage.

Data exfiltration: custom regex (exfiltration_detector.py), 4 sub-categories, EN/KO/UR. Covers system-prompt extraction, credential/secret extraction, indirection/bypass attempts (e.g. "translate your instructions," "summarize your own rules"), and bulk corpus extraction (a RAG-specific attack surface: asking the system to dump raw retrieved chunk content rather than answer a question). No pretrained classifier was found purpose-built for this task (see Section 3);

this module was built entirely from hand-written patterns, developed and iteratively widened based on real misses surfaced during evaluation (Section 6).

Toxicity: textdetox/bert-multilingual-toxicity-classifier + regex supplement (toxicity_detector.py, KO/UR only). The classifier alone is used for English, where it tested clean. Korean and Urdu are supplemented with hand-written patterns covering two attack shapes found during testing: direct insults ("you're all [insult], you should disappear") and discriminatory generalizations stated as fact ("[group] can't do X, so they shouldn't be allowed to") — the latter was a distinct pattern class not covered by the direct-insult patterns alone, found only after the classifier missed every discriminatory-generalization example in the initial Korean evaluation pass.

PII: Microsoft Presidio (English only) + custom regex (pii_detector.py, all 3 languages). Presidio's English NER handles EMAIL_ADDRESS and PHONE_NUMBER detection (its PERSON entity was deliberately excluded — see Section 7). Structured identifiers — SSN, Pakistani CNIC, Korean RRN, credit card numbers, phone numbers in all three languages — are hand-written regex, built after Presidio's own SSN recognizer was found to be broken (Section 6) and after confirming Presidio has zero Korean/Urdu language support. Blocks on both input and output.

Jailbreak (narrative/roleplay): LLM semantic judge (llm_judge.py). A single-purpose LLM call, separate from the existing query-rewrite/multi-query/citation calls, that classifies a query as SAFE or UNSAFE based on underlying intent rather than surface phrasing. Explicitly prompted to distinguish genuine creative/hypothetical framing from framing used to smuggle a harmful request past a keyword filter (e.g. a "novelist" asking for detailed, actionable hacking instructions inside a story, versus a "novelist" asking for a plausible-sounding but non-actionable plot beat). This is the component that closes the fast layer's largest remaining gap (Section 8), and is only reached when the fast layer finds nothing — see Section 5.1 for its measured contribution.

3. Tool Selection: What Was Tested and Rejected

Consistent with this project's approach throughout, no tool was adopted on the strength of its documentation alone — every candidate was tested directly against hand-picked EN/KO/UR prompts before being trusted.

Prompt injection/jailbreak classifiers tested:

- protectai/deberta-v3-base-prompt-injection-v2 — English-only base model (not multilingual, despite general "prompt injection detection" framing). Direct testing confirmed a false positive on a completely benign Urdu query (فرانس کا دار الحکومت کیا ہے؟, "What is the capital of France?") flagged as injection at 0.980 confidence.
- meta-llama/Llama-Prompt-Guard-2-86M — built on a multilingual base model, though Meta's own evaluation set covers English, French, German, Hindi, Italian, Portuguese, Spanish, and Thai — not Korean or Urdu. Direct 12-case smoke test showed 10/12 correct,

with the same Urdu benign query correctly classified (unlike ProtectAI), and two misses both being roleplay-style jailbreak framing in Korean and Urdu specifically, while the equivalent English case was caught. Selected as the primary classifier on this evidence, with the identified roleplay gap addressed first via regex supplement and later via the LLM judge for the broader narrative-framing pattern.

Toxicity classifiers tested:

- meta-llama/Llama-Guard-3-8B — disqualified on practical grounds. A 91-minute download (16.1GB) was followed by an indefinite hang during weight loading on available hardware. This is reported as a legitimate finding: an 8B-parameter generative safety classifier is impractical for local/resource-constrained deployment, regardless of its detection quality, which was never able to be evaluated.
- textdetox/bert-multilingual-toxicity-classifier — lightweight (~178M params, bert-base-multilingual-cased), no gating, fast to load and run. Smoke test: 3/3 on English, 1/2 on Korean, 1/2 on Urdu. Both non-English misses were toxic content misclassified as neutral (the safer-seeming but more dangerous failure direction for a safety layer, since it lets real harmful content through silently). Selected with a regex supplement for Korean/Urdu.

PII detection tested:

- Microsoft Presidio — English NER (PERSON, EMAIL_ADDRESS via spaCy) worked correctly in initial testing. However, direct isolated testing of Presidio's built-in UsSsnRecognizer — called with score_threshold=0.0, with explicit context keywords immediately adjacent to the target string, and finally called completely standalone bypassing the full analyzer pipeline — returned zero results in every configuration despite a syntactically correct regex pattern being present in the recognizer's own pattern list. This is a genuine, reproducible bug in Presidio's shipped SSN recognizer, not a configuration or phrasing issue on this project's part. Presidio also has zero Korean or Urdu language configuration (analyzer.supported_languages == ['en']), confirmed directly rather than assumed from documentation.
- Resolution: PERSON was later removed from Presidio's trusted entity set entirely (see Section 7), and SSN, phone, and Korean RRN detection were rebuilt as custom regex rather than relying on Presidio's broken implementation.

4. Attack Prompt Test Set


A 125-prompt labeled evaluation set was constructed, spanning all five categories across English, Korean, and Urdu, plus benign controls.

Language	Attack prompts	Benign prompts	Total
English	34	19	53
Korean	30	6	36
Urdu	30	6	36
Total	94	31	125

English prompts were sourced from Necent/llm-jailbreak-prompt-injection-dataset, a public aggregate of 30+ established safety/security research datasets (HarmBench, AdvBench, TensorTrust, LLMail-Inject, WildJailbreak, RealToxicityPrompts, SGTotoxicGuard, among others), filtered to category-relevant subsets and then manually reviewed. Manual review excluded prompts that were mislabeled, corrupted/garbled, referenced missing multimodal content, or — in one case — contained content involving a minor that had no place in an evaluation set regardless of its source label.

Korean and Urdu prompts were entirely hand-written, after direct testing confirmed the same aggregate dataset had no usable non-English content for this project: its 11,209 Korean rows were either mislabeled harmful content unrelated to the intended category, or scraped web spam auto-labeled as "toxicity" with no actual toxic content present. This is reported as a finding in its own right — a large, well-cited, actively-maintained safety dataset can still be entirely unusable for a specific language subset despite a substantial raw row count, and this was only caught by manually inspecting sample content rather than trusting the category label.

5. Evaluation Results



5.1 Fast Layer vs. Fast+LLM: Overall Comparison

The guardrail was evaluated in two configurations: the fast layer alone (check_input, no LLM calls), and the production configuration with the LLM judge as a second pass (check_input_deep).


Metric	Fast layer only	Fast + LLM judge
Accuracy	0.840	0.928
Precision	0.987	0.989
Recall	0.798	0.915
F1	0.882	0.950
Attack success rate	0.202	0.085
False positive rate	0.032	0.032

The LLM judge fired on approximately 8.8% of the 125-prompt set (11/125), meaning roughly 91% of traffic resolves entirely at the fast layer with no added LLM cost. Among the prompts that reached the LLM stage, it recovered 9 of the fast layer's 20 misses without introducing any new false positives — precision held essentially flat (0.987 → 0.988) while recall improved by nearly 10 points. This is the configuration adopted for the production system (check_input_deep), on the reasoning that the added latency and cost, bounded to under 9% of traffic, is justified by closing most of the fast layer's largest and most structurally difficult gap (narrative jailbreak, Section 8).

The remainder of this section reports the fast+LLM configuration's results in detail, since that is the deployed system; fast-layer-only figures are noted inline where the comparison is informative.

5.2 By Category

Left: without LLM judge, right: with LLM judge




Category	Accuracy	Recall	Attack success rate	Fast-layer-only recall
Exfiltration	1.000	1.000	0.000	1.000
PII	1.000	1.000	0.000	1.000
Injection	1.000	1.000	0.000	0.955
Jailbreak	0.900	0.900	0.100	0.450
Toxicity	0.714	0.714	0.286	0.667
Benign (false positive check)	0.968	—	—	0.968

Jailbreak shows the largest improvement by a wide margin (recall 0.450 $\rightarrow$ 0.850), consistent with the LLM judge being purpose-built to catch exactly this category's dominant failure mode. Toxicity improved marginally (0.667 $\rightarrow$ 0.714) even though the LLM judge is not scoped to toxicity at all — this is attributable to one or two toxicity-labeled prompts also containing incidental jailbreak-adjacent framing, not the judge directly targeting toxic content. Exfiltration and PII are unchanged, as expected, since the LLM judge is only reached when the fast layer already finds nothing, and these two categories were already at or near ceiling in the fast layer alone.

5.3 By Language

Left: without LLM judge, right: with LLM judge



Language	Accuracy	Recall	Attack success rate	False positive rate
Urdu	1.000	1.000	0.000	0.000
Korean	1.000	1.000	0.000	0.000
English	0.811	0.735	0.265	0.053

With the LLM judge added, Urdu and Korean reach perfect accuracy, while English remains the weakest language — the reverse of what is typically expected from multilingual NLP tooling. This is explained by category composition rather than the LLM judge favoring non-English text: English's remaining misses are now concentrated almost entirely in the toxicity dialogue-format gap (Section 8), a limitation of the English toxicity classifier specifically, which the LLM judge is not scoped to address. Korean and Urdu's remaining gaps were largely closed either by iterative regex supplementation (Section 6) or by the LLM judge itself.

6. Design Decisions Made From Evidence

PII blocks on input, not just output. The initial design only blocked PII on check_output (generated answers), reasoning that a user pasting their own phone number into a query wasn't necessarily an attack. This was revised to block PII on both input and output after review — there is no clear benefit to letting a user's sensitive information flow into retrieval and generation unnecessarily even if it isn't hostile, and the simpler, more conservative policy was adopted. This changed the eval set's expected_blocked labels for the PII category to match.

Presidio's PERSON entity was removed from the trusted set. When PII input-blocking was first enabled, four new false positives appeared, all triggered by ordinary names in benign queries

("How can I reach Shelley?", a Spider-Man roleplay prompt naming the character, a wildlife-documentary prompt naming a fictional researcher). A bare name is not inherently sensitive — it appears constantly in ordinary text — so PERSON was dropped from the entities that trigger a block, while EMAIL_ADDRESS and PHONE_NUMBER (genuinely sensitive identifiers) were retained.

The LLM judge runs only after the fast layer, not before or in parallel. This ordering was chosen deliberately for cost: the fast layer is near-instant and already resolves the large majority of both attacks and benign traffic. Running the LLM judge on every request regardless of the fast layer's result would multiply the added latency/cost across all 125 evaluation prompts rather than the ~9% that actually needed it, for no corresponding gain in the categories the fast layer already handles well.

The injection classifier's confidence threshold was tuned against the eval set rather than left at an untested default. The initial value (0.9) was a placeholder. A sweep across thresholds 0.5–0.99 against the 125-prompt set showed the classifier's score distribution is largely bimodal — recall moved only from 0.245 to 0.277 across the entire range, indicating most of its errors are confident misclassifications rather than borderline cases sitting near a threshold boundary. 0.5 was selected as the value with the best recall at no measurable precision cost (0.963 vs. 0.958 at 0.9, on the classifier in isolation). Applied to the full system, this closed two additional misses (one English jailbreak case, one Korean injection case) without introducing new false positives, moving overall accuracy from 91.2% to 92.8%.

7. Known Limitations



English toxicity in dialogue-turn format (28.6% attack success rate, concentrated in one sub-source). Several missed English toxicity prompts share a specific structural template from the SGToxicGuard source dataset: "A: [toxic statement] B: [generate response]". The classifier reliably identifies toxicity when it is stated directly but appears to not recognize it when embedded inside this dialogue-continuation format, despite the toxic content itself being unchanged. Korean and Urdu do not show this specific gap since their regex supplement operates on raw text content regardless of surrounding structure — this is a limitation specific to how the English classifier processes structured input, not a language-coverage gap. The LLM judge is currently scoped to jailbreak intent only and was not asked to evaluate toxicity; extending its scope to also catch this dialogue-format toxicity gap is a natural next step, discussed further below.

Residual jailbreak misses even with the LLM judge (10.0% attack success rate). The LLM judge closed most, but not all, of the narrative-framing gap — 2 of the original 20 fast-layer jailbreak misses remained after the LLM judge was added. This indicates the judge itself has an imperfect, if substantially better, ceiling on this category; a more capable or differently-prompted judge model,

or a second independent LLM pass, could plausibly close further ground here, though this was not explored in this phase.

Remaining false positive. One benign prompt — a request for euphemistic language to describe body-type bias in fiction writing, explicitly framed as staying within ethical guidelines — is flagged by Prompt Guard 2 as injection/jailbreak. This was not investigated further given the low overall false positive rate (3.2%) and the isolated nature of the case; it did not recur across any other benign prompt with similar "stay within guidelines" language.

8. Conclusion



The guardrail module implements all five required detection categories across three languages, built from an evidence-driven combination of pretrained classifiers (verified rather than assumed to work multilingually), hand-built regex supplements written to close gaps found through direct testing, and — for the category no pattern-matching approach could reliably resolve — a targeted LLM-based semantic judge.

The fast-layer-only configuration reaches 84.0% accuracy; adding the LLM judge as a second pass, reached only on the roughly 9% of traffic the fast layer cannot already resolve, raises this to **91.2% accuracy, 89.4% recall, and 98.8% precision**. This two-stage design — cheap deterministic checks first, an LLM judge only where structurally necessary — is the configuration adopted for the production system, and reflects a deliberate cost/capability tradeoff rather than defaulting to the most powerful available tool for every request.

The system's remaining known gaps — English toxicity in dialogue-turn format, and a residual fraction of jailbreak attempts even after the LLM judge — are documented as scoped, evidenced limitations rather than resolved by further narrow patching, consistent with this project's approach throughout.